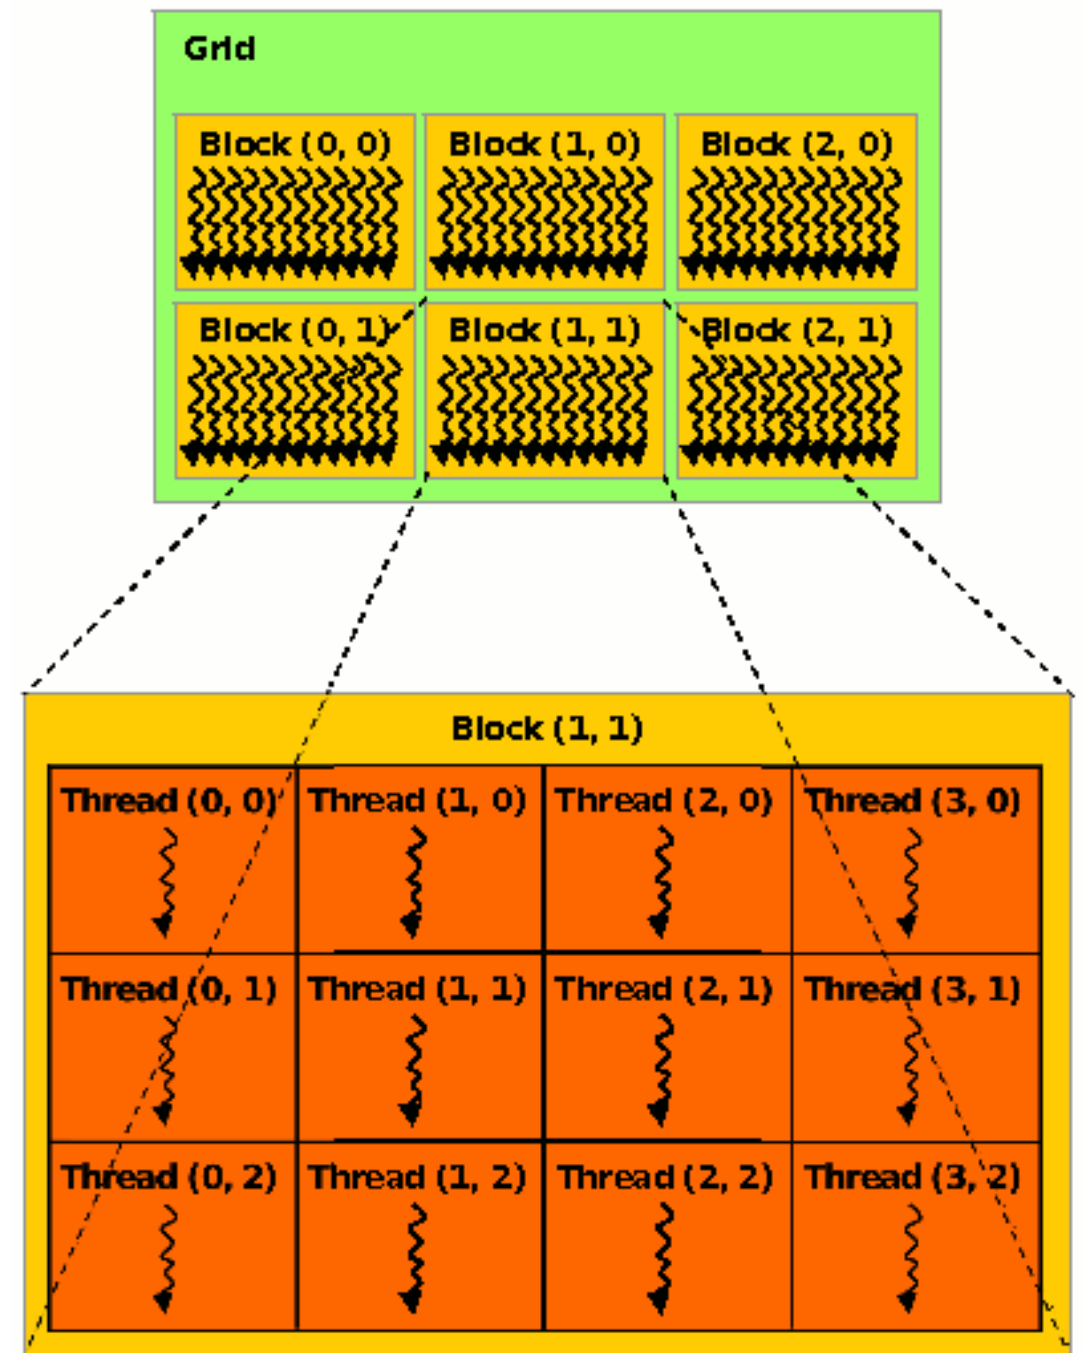



# NVIDIA GPU Memory Hierarchy

- Grids map to GPUs
- Blocks map to the MultiProcessors (MP)
- Threads map to Stream Processors (SP)
- Warps are groups of (32) threads that execute simultaneously



# NVIDIA GPU Memory Architecture

- In a NVIDIA GTX 480:
  - Maximum number of threads per block:  
1024
  - Maximum sizes of x-, y-, and z- dimensions of thread block:  
1024 x 1024 x 64
  - Maximum size of each dimension of grid of thread blocks:  
65535 x 65535 x 65535



# Defining Grid/Block Structure

- Need to provide each kernel call with values for two key structures:
  - Number of blocks in each dimension
  - Threads per block in each dimension
- `myKernel<<< B,T >>>(arg1, ... );`
- B – a structure that defines the number of blocks in grid in each dimension (1D or 2D).
- T – a structure that defines the number of threads in a block in each dimension (1D, 2D, or 3D).

# 1D Grids and/or 1D Blocks

- If want a 1-D structure, can use a integer for B and T in:
- `myKernel<<< B, T >>>(arg1, ... );`
- B – An integer would define a 1D grid of that size
- T – An integer would define a 1D block of that size
- Example: `myKernel<<< 1, 100 >>>(arg1, ... );`

# CUDA Built-In Variables

- **blockIdx.x, blockIdx.y, blockIdx.z** are built-in variables that return the block ID in the x-axis, y-axis, and z-axis of the block that is executing the given block of code.
- **threadIdx.x, threadIdx.y, threadIdx.z** are built-in variables that return the thread ID in the x-axis, y-axis, and z-axis of the thread that is being executed by this stream processor in this particular block.
- **blockDim.x, blockDim.y, blockDim.z** are built-in variables that return the “block dimension” (i.e., the number of threads in a block in the x-axis, y-axis, and z-axis).
- So, you can express your collection of blocks, and your collection of threads within a block, as a 1D array, a 2D array or a 3D array.
- These can be helpful when thinking of your data as 2D or 3D.
- The full global thread ID in x dimension can be computed by:
  - $x = \text{blockIdx.x} * \text{blockDim.x} + \text{threadIdx.x};$

# Thread Identification Example: x-direction

Global Thread ID

|                |   |   |   |   |   |   |   |                |   |    |    |    |    |    |    |                |    |    |    |    |    |    |    |                |    |    |    |    |    |    |    |
|----------------|---|---|---|---|---|---|---|----------------|---|----|----|----|----|----|----|----------------|----|----|----|----|----|----|----|----------------|----|----|----|----|----|----|----|
| 0              | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8              | 9 | 10 | 11 | 12 | 13 | 14 | 15 | 16             | 17 | 18 | 19 | 20 | 21 | 22 | 23 | 24             | 25 | 26 | 27 | 28 | 29 | 30 | 31 |
| threadIdx.x    |   |   |   |   |   |   |   | threadIdx.x    |   |    |    |    |    |    |    | threadIdx.x    |    |    |    |    |    |    |    | threadIdx.x    |    |    |    |    |    |    |    |
| 0              | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 0              | 1 | 2  | 3  | 4  | 5  | 6  | 7  | 0              | 1  | 2  | 3  | 4  | 5  | 6  | 7  | 0              | 1  | 2  | 3  | 4  | 5  | 6  | 7  |
| blockIdx.x = 0 |   |   |   |   |   |   |   | blockIdx.x = 1 |   |    |    |    |    |    |    | blockIdx.x = 2 |    |    |    |    |    |    |    | blockIdx.x = 3 |    |    |    |    |    |    |    |

- Assume a hypothetical ID grid and ID block architecture: 4 blocks, each with 8 threads.
- For Global Thread ID 26:
  - $\text{gridDim.x} = 4 \times 1$
  - $\text{blockDim.x} = 8 \times 1$
  - $\text{Global Thread ID} = \text{blockIdx.x} * \text{blockDim.x} + \text{threadIdx.x}$
  - $= 3 \times 8 + 2 = 26$